Google Cloud Alphafold Server Statistics

Jason Zhu

Single Process Performance

Tesla T4 $255.50


```
{
  "features": 1236.4132311344147,
  "process_features_model_1_pred_0": 4.506194591522217,
  "predict_and_compile_model_1_pred_0": 123.11969947814941,
  "relax_model_1_pred_0": 22.15885877609253,
```


Tesla K80 $328.50


```
{
  "features": 1544.4125242233276,
  "process_features_model_1_pred_0": 5.009815692901611,
  "predict_and_compile_model_1_pred_0": 157.80067491531372,
  "relax_model_1_pred_0": 21.56758165359497,
```


Tesla P4 $438.00


```
{
  "features": 1044.6379506587982,
  "process_features_model_1_pred_0": 4.024508714675903,
  "predict_and_compile_model_1_pred_0": 112.10066890716553,
  "relax_model_1_pred_0": 19.794049739837646,
```

Single Process Performance

Tesla P100 $1065.80


```
{
    "features": 1286.3185884952545,
    "process_features_model_1_pred_0": 4.223628282546997,
    "predict_and_compile_model_1_pred_0": 117.43888092041016,
    "relax_model_1_pred_0": 22.24559473991394,
```


Tesla V100 $1810.40


```
{
    "features": 1442.019693851471,
    "process_features_model_1_pred_0": 4.730117559432983,
    "predict_and_compile_model_1_pred_0": 125.50818347930908,
    "relax_model_1_pred_0": 22.482089996337 89,
```


Tesla A100 $2141.75


```
{
    "features": 1121.5335009098053,
    "process_features_model_1_pred_0": 3.867100477218628,
    "predict_and_compile_model_1_pred_0": 225.68341636657715,
    "relax_model_1_pred_0": 20.6371679306030 27,
```

Multiple Process (3 processes) Performance

Tesla T4 $255.50


```
{
    "features": 1433.34395098686 22,
    "process_features_model_1_pred_0": 4.936453580856323,
    "predict_and_compile_model_1_pred_0": 170.94763731956482,
    "relax_model_1_pred_0": 25.750863313674927,
```


Tesla K80 $328.50


```
{
    "features": 1952.3416156768799,
    "process_features_model_1_pred_0": 5.5588972568511 96,
    "predict_and_compile_model_1_pred_0": 159.50953674316406,
    "relax_model_1_pred_0": 22.533478021621704,
```


Tesla P4 $438.00


```
{
    "features": 1449.0862200260162,
    "process_features_model_1_pred_0": 3.7654290199279785,
    "predict_and_compile_model_1_pred_0": 106.27897429466248,
    "relax_model_1_pred_0": 18.74294686317444,
```

Multiple Process (3 processes) Performance

Tesla P100 $1065.80


```
{
    "features": 1438.4974267482758,
    "process_features_model_1_pred_0": 4.949854850769043,
    "predict_and_compile_model_1_pred_0": 164.7899477481842,
    "relax_model_1_pred_0": 23.740139722824097,
```


Tesla V100 $1810.40


```
{
    "features": 1638.1646158695 22,
    "process_features_model_1_pred_0": 5.0536651611328125,
    "predict_and_compile_model_1_pred_0": 170.21040773391724,
    "relax_model_1_pred_0": 26.750015020370483,
```


Tesla A100 $2141.75


```
{
    "features": 1405.3440382480621,
    "process_features_model_1_pred_0": 4.513792514801025,
    "predict_and_compile_model_1_pred_0": 298.0864064693451,
    "relax_model_1_pred_0": 21.606945276260376,
```

General Analysis

- there’s no clear relationship between gpu prices and running time
- one process takes maximum of 225% vCPU resource each process, so for setting of 8 vCPUs & 52GB RAM, it supports 3 simultaneous processes.
- gpu can help to boost the overall speed, while the effect of multiple gpu is not yet clear
- adding more CPU & correponding RAM would increase the maximum simultaneous process - $33 for each vCPU&RAM on google cloud

Recommandation Based on Current Data

- 3TB SSD disk for data storage - $307.2 monthly
- 8 vCPU & 52GB RAM ($345.44) with 1 teslaT4 GPU ($255.50) as a basic set
- 100 GB boot disk ($10 monthly)
- 725 seconds for one GFP11 mutation prediction (16 amino acids)
- to increase running speed, only need to add more basic set with no need for SSD and boot disk

Item	Monthly estimate
8 vCPU + 52 GB memory	$345.44
1 nvidia-tesla-t4	$255.50
100 GB balanced persistent disk	$10.00
Sustained use discount	-$180.28
Total	$430.66

shell script for auto-initialization

```
sudo mkdir -p /mnt/disks/data
sudo mount -o discard,defaults /dev/sdb /mnt/disks/data
sudo chmod a+w /mnt/disks/data
sudo apt-get remove docker docker-engine docker.io containerd runc
sudo apt-get update
sudo apt-get install \
    ca-certificates \
    curl \
    gnupg \
    lsb-release
sudo mkdir -p /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/debian/gpg | sudo gpg --dearmor -o /etc/apt/keyrings
echo \
      "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://dow
        $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update
sudo apt-get install docker-ce docker-ce-cli containerd.io docker-compose-plugin
#sudo groupadd docker
sudo usermod -aG docker $USER
newgrp docker
curl https://raw.githubusercontent.com/GoogleCloudPlatform/compute-gpu-installation/main/linux/in
sudo python3 install_gpu_driver.py
distribution=$(. /etc/os-release;echo $ID$VERSION_ID) \
    && curl -fsSL https://nvidia.github.io/libnvidia-container/gpgkey | sudo gpg --dearmor -o /us
    && curl -s -L https://nvidia.github.io/libnvidia-container/$distribution/libnvidia-container.
    sed 's#deb https://#deb [signed-by=/usr/share/keyrings/nvidia-container-toolkit-keyring.gpg]
    sudo tee /etc/apt/sources.list.d/nvidia-container-toolkit.list
sudo apt-get update
sudo apt-get install -y nvidia-docker2
sudo systemctl restart docker
sudo docker run --rm --gpus all nvidia/cuda:11.0.3-base-ubuntu20.04 nvidia-smi
sudo apt-get install git
git clone https://github.com/deepmind/alphafold.git
```

Shell script for one step running

```
#!/bin/bash

if [ $# -lt 1 ]; then
    echo "Usage: ./run_alphafold.sh [input fasta path] [output directory]"
    exit 1
fi
input=$1
if [ $# -eq 1 ]; then
    output="/tmp/alphafold"
else
    output=$2
fi

if [ -d $output ]; then
    echo "saving results to ${output}"
else
    mkdir $output
    echo "saving results to ${output}"
fi

python3 docker/run_docker.py --fasta_paths=$input \
    --max_template_date=2021-05-14 --output_dir=$output \
    --data_dir=/mnt/disks/data
```

Python script for auto-mutation generation

alphafold_user@alphafold-server-gpu-p4-0-99: ~/alphafold/seq_input

```
    while curr_line != "":
        curr_line = f.readline().strip()
        body += curr_line

#print(heading)
#print(body)
if argc == 2 and sys.argv[1] == "all":
    start = 0
    end = len(body)
elif argc == 2 and sys.argv[1].isdigit():
    start = int(sys.argv[1])
    end = start + 1
elif argc == 3 and sys.argv[1].isdigit() and sys.argv[2].isdigit():
    start = int(sys.argv[1])
    end = int(sys.argv[2])
    if start < 0 or start >= end or end > len(body):
        print("Invalid Interval input")
        exit()
else:
    print("Invalid non-digit input")
    exit()
#print(f"start:{start}, end:{end}")

for i in range(start, end):
    for j in range(len(amino_seq)):
        if body[i] == amino_seq[j]:
            continue
        name = f"mutate_amino{i}_to_{amino_seq[j]}.fasta"
        #print(name)
        with open(name, 'w') as f:
            print(heading, file = f)
            print(body[0: i] + amino_seq[j] + body[i + 1:], file = f)


"generate_mutation.py" 48L, 1489B
```

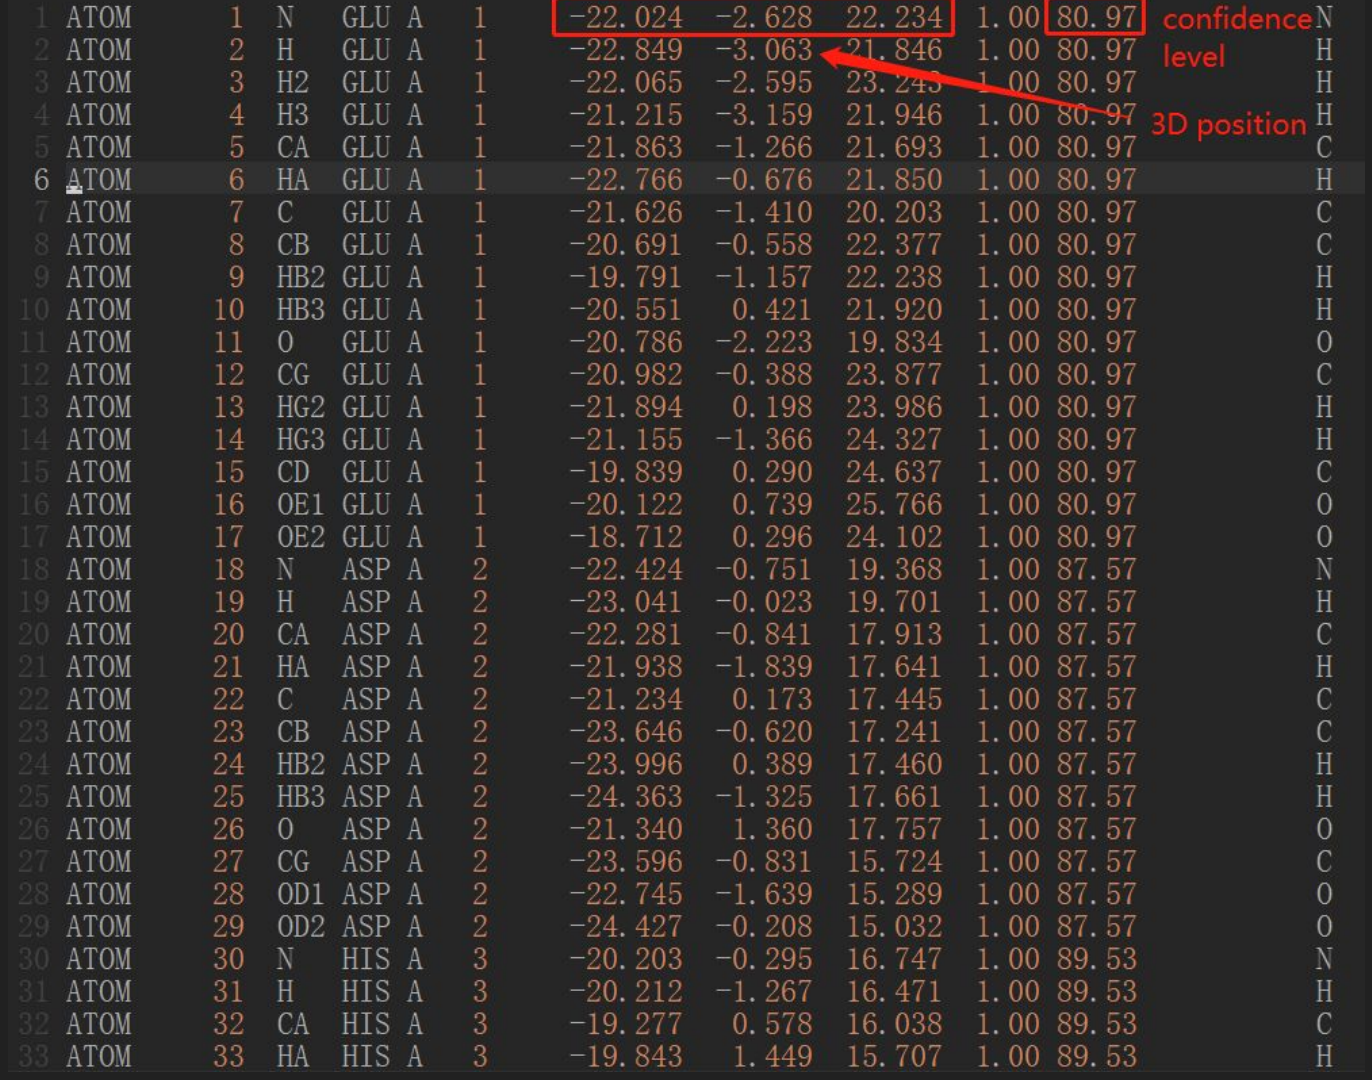

1	ATOM	1	N	GLU	A	1	-22.024	-2.628	22.234	1.00	80.97	N
2	ATOM	2	H	GLU	A	1	-22.849	-3.063	21.846	1.00	80.97	H
3	ATOM	3	H2	GLU	A	1	-22.065	-2.595	23.245	1.00	80.97	H
4	ATOM	4	H3	GLU	A	1	-21.215	-3.159	21.946	1.00	80.97	H
5	ATOM	5	CA	GLU	A	1	-21.863	-1.266	21.693	1.00	80.97	C
6	ATOM	6	HA	GLU	A	1	-22.766	-0.676	21.850	1.00	80.97	H
7	ATOM	7	C	GLU	A	1	-21.626	-1.410	20.203	1.00	80.97	C
8	ATOM	8	CB	GLU	A	1	-20.691	-0.558	22.377	1.00	80.97	C
9	ATOM	9	HB2	GLU	A	1	-19.791	-1.157	22.238	1.00	80.97	H
10	ATOM	10	HB3	GLU	A	1	-20.551	0.421	21.920	1.00	80.97	H
11	ATOM	11	O	GLU	A	1	-20.786	-2.223	19.834	1.00	80.97	O
12	ATOM	12	CG	GLU	A	1	-20.982	-0.388	23.877	1.00	80.97	C
13	ATOM	13	HG2	GLU	A	1	-21.894	0.198	23.986	1.00	80.97	H
14	ATOM	14	HG3	GLU	A	1	-21.155	-1.366	24.327	1.00	80.97	H
15	ATOM	15	CD	GLU	A	1	-19.839	0.290	24.637	1.00	80.97	C
16	ATOM	16	OE1	GLU	A	1	-20.122	0.739	25.766	1.00	80.97	O
17	ATOM	17	OE2	GLU	A	1	-18.712	0.296	24.102	1.00	80.97	O
18	ATOM	18	N	ASP	A	2	-22.424	-0.751	19.368	1.00	87.57	N
19	ATOM	19	H	ASP	A	2	-23.041	-0.023	19.701	1.00	87.57	H
20	ATOM	20	CA	ASP	A	2	-22.281	-0.841	17.913	1.00	87.57	C
21	ATOM	21	HA	ASP	A	2	-21.938	-1.839	17.641	1.00	87.57	H
22	ATOM	22	C	ASP	A	2	-21.234	0.173	17.445	1.00	87.57	C
23	ATOM	23	CB	ASP	A	2	-23.646	-0.620	17.241	1.00	87.57	C
24	ATOM	24	HB2	ASP	A	2	-23.996	0.389	17.460	1.00	87.57	H
25	ATOM	25	HB3	ASP	A	2	-24.363	-1.325	17.661	1.00	87.57	H
26	ATOM	26	O	ASP	A	2	-21.340	1.360	17.757	1.00	87.57	O
27	ATOM	27	CG	ASP	A	2	-23.596	-0.831	15.724	1.00	87.57	C
28	ATOM	28	OD1	ASP	A	2	-22.745	-1.639	15.289	1.00	87.57	O
29	ATOM	29	OD2	ASP	A	2	-24.427	-0.208	15.032	1.00	87.57	O
30	ATOM	30	N	HIS	A	3	-20.203	-0.295	16.747	1.00	89.53	N
31	ATOM	31	H	HIS	A	3	-20.212	-1.267	16.471	1.00	89.53	H
32	ATOM	32	CA	HIS	A	3	-19.277	0.578	16.038	1.00	89.53	C
33	ATOM	33	HA	HIS	A	3	-19.843	1.449	15.707	1.00	89.53	H

shell/python program for automatic running

- ssh/mobaxterm for login, all have linux interface avaliable
- scp for transferring data back / store in server locally